

BEYOND PERMUTATION SYMMETRY: WIDTH AND SCALING IN LINEAR MODE CONNECTIVITY

Anonymous authors

Paper under double-blind review

ABSTRACT

Permutation symmetry is often used to explain why independently trained neural networks appear separated under linear interpolation. We test where that explanation is exact, where it depends on width, and whether a larger symmetry class is needed. On XOR, exhaustive enumeration gives an oracle over all hidden-unit permutations. Minimal-width ReLU networks provide concrete pairs for which every permutation retains a loss barrier, ruling out permutation search failure in this setting. Increasing width rapidly lowers the oracle barrier, while a separate nonlinear-path experiment exhibits a sharp change between widths two and three, separating symmetry alignment from representational slack. We then exploit positive scaling symmetry and propose a scale-aware extension of differentiable Sinkhorn re-basing. Starting from the oracle permutation and optimizing only scales reduces the mean XOR barrier from 0.357 to 0.247, 0.111 to 0.025, and 0.023 to 0.005 at hidden widths 3, 5, and 7, respectively, and never worsens a pair. A shared-training experiment shows that permutation alignment becomes harder as training trajectories diverge, and experiments on VGG networks trained on CIFAR-10 show that adding scale improves over permutation-only alignment in all four tested architectures. These results refine the one-basin picture: permutation explains part of linear compatibility, width supplies additional slack, and positive scaling is a practically relevant exact symmetry beyond permutation.

1 INTRODUCTION

Neural networks trained from different random initializations usually occupy distant points in parameter space, and the straight line between their weights often crosses a region of high loss. This observation appears to suggest that stochastic gradient descent (SGD) finds distinct basins. Yet neural-network parameterizations are highly non-identifiable. Reordering hidden units or channels, with the inverse reordering applied to the next layer, preserves the represented function. Consequently, a barrier between two raw parameter vectors may reflect an arbitrary choice of representative rather than genuine geometric separation.

The one-basin modulo permutation hypothesis makes this intuition precise: most SGD solutions may admit functionally equivalent permutations that are linearly mode connected (Entezari et al., 2021). Explicit alignment methods have made this hypothesis operational. Git Re-Basin substantially reduces barriers for MLPs and convolutional networks and obtains essentially barrier-free interpolation in several sufficiently wide settings (Ainsworth et al., 2023). However, residual barriers remain for more difficult architectures and smaller widths. Such failures have two incompatible explanations: the search method may miss a good permutation, or no good permutation may exist. Modern networks have combinatorially large permutation groups, so standard experiments cannot separate these possibilities.

We use a small setting in which the ambiguity disappears. For one-hidden-layer XOR networks, every permutation can be evaluated exactly. This controlled test reveals trained pairs for which permutation is genuinely insufficient, and a width sweep shows how quickly the failure weakens as redundant units are added. We complement the linear test with nonlinear path optimization at widths two and three. This distinguishes two mechanisms that can otherwise be conflated: selecting a better representative using exact symmetries, and using extra model capacity to move through a broader low-loss set.

The failure of an exhaustive permutation search motivates a larger exact symmetry class. ReLU networks are invariant to positive rescaling of a hidden unit when the outgoing weight is inversely rescaled. We incorporate these continuous degrees of freedom into Sinkhorn re-basing (Peña et al., 2023) and also study a cleaner oracle protocol: first choose the best exhaustive permutation, then optimize scales while freezing that permutation. The latter isolates the benefit of scaling from improvements in permutation search.

Our contributions are:

- An exhaustive XOR counterexample showing that nonzero barriers can remain after the best possible permutation, together with a width sweep demonstrating that the oracle barrier collapses as redundancy grows.
- A nonlinear-connectivity comparison between widths two and three, showing that one extra hidden unit can change path connectivity even when permutation-based LMC remains imperfect.
- A scale-aware Sinkhorn parameterization and an oracle permutation-plus-scale experiment. Scale refinement improves or preserves the barrier for all 67 tested XOR pairs and produces the largest gains at widths five and seven.
- A training-stage analysis on VGG16 and a realistic validation on VGG11/13/16/19 trained on CIFAR-10. The former links alignment difficulty to training divergence; the latter shows consistent gains from scale beyond permutation-only alignment.

2 BACKGROUND AND RELATED WORK

Connectivity and interpolation barriers. For endpoints θ_A and θ_B , linear interpolation is $\theta(t) = (1 - t)\theta_A + t\theta_B$. Following prior work, for a scalar criterion C (loss or classification error) we measure the excess above the linear interpolation of endpoint values:

$$\mathcal{B}(\theta_A, \theta_B; C) = \max_{t \in [0,1]} [C(\theta(t)) - (1 - t)C(\theta_A) - tC(\theta_B)]. \quad (1)$$

The maximum is approximated on a dense grid. Low-barrier curved paths between independent solutions are widespread (Draxler et al., 2018; Garipov et al., 2018), but curved connectivity is weaker than linear mode connectivity (LMC): the path may avoid a region crossed by the straight segment. Shared-initialization experiments further show that solutions become linearly stable to SGD noise early in training (Frankle et al., 2020).

Why linear connectivity matters. A curved connection establishes geometric connectedness, but a straight path is the relevant object for weight averaging. Averaging checkpoints from one training trajectory can improve generalization by favoring a wider region of parameter space (Izmailov et al., 2019), and averaging related fine-tuned models can outperform the individual endpoints without ensemble-time cost (Wortsman et al., 2022). Independent models are harder: if their hidden features use incompatible coordinates, the midpoint mixes unrelated units and can be unusable even when both endpoints are accurate. Re-basing seeks a functionally equivalent coordinate system before averaging. Our results focus on the prerequisite—a low-barrier segment—rather than claiming that the midpoint necessarily improves endpoint accuracy. This distinction matters because a symmetry transformation can repair compatibility while providing no new predictive information by itself.

Permutation re-basing. For a ReLU MLP block $f(x) = W_2\sigma(W_1x)$ and a permutation matrix P ,

$$(W_2P^\top)\sigma(PW_1x) = W_2\sigma(W_1x). \quad (2)$$

The same action reorders channels in convolutional networks. Weight matching finds layerwise permutations that reduce Euclidean distance, activation matching aligns representations on data, and differentiable methods optimize soft assignments (Ainsworth et al., 2023; Peña et al., 2023). REPAIR shows that correcting activation statistics can further improve interpolation after permutation (Jordan et al., 2023). These successes support a symmetry-aware view of loss landscapes, but a residual barrier obtained by a heuristic cannot establish that the symmetry class itself is insufficient.

Positive scaling symmetry. Positive homogeneity, $\sigma(cz) = c\sigma(z)$ for $c > 0$, yields a second exact symmetry. With D positive diagonal,

$$(W_2 D^{-1})\sigma(DW_1 x) = W_2 \sigma(W_1 x). \quad (3)$$

Permutation and scale therefore form a monomial action $M = PD$. Unlike a permutation, scaling moves a representative continuously and can substantially change its Euclidean relation to a second model. Our question is not whether scaling changes either endpoint function—it does not—but whether it selects representatives whose straight-line interpolation is better behaved.

3 SCALE-AWARE ALIGNMENT

Sinkhorn re-basing replaces a hard layerwise permutation by a differentiable, approximately doubly stochastic matrix

$$S_\ell = \text{Sinkhorn}(Z_\ell/\tau), \quad (4)$$

where Z_ℓ is an unconstrained score matrix and τ is a temperature. Repeated row and column normalization of $\exp(Z_\ell/\tau)$ permits gradient optimization through the alignment. After optimization, the soft matrix is projected to a hard permutation using linear assignment.

We augment each hidden layer with learnable log-scales u_ℓ :

$$D_\ell = \text{diag}(e^{u_\ell}), \quad M_\ell = S_\ell D_\ell. \quad (5)$$

On the output coordinates of layer ℓ we apply $S_\ell D_\ell$; on the input coordinates of the next layer we apply the compensating $D_\ell^{-1} S_\ell^\top$. Once S_ℓ is projected to a permutation, this is the exact function-preserving monomial transformation in Eqs. 2–3. During soft optimization, the transpose is the usual differentiable Sinkhorn approximation rather than an exact inverse.

The variables are optimized on real data using midpoint interpolation loss:

$$J = \mathcal{L}(\tfrac{1}{2}\theta_A + \tfrac{1}{2}M(\theta_B)) + \lambda \sum_\ell \|u_\ell\|_2^2. \quad (6)$$

The regularizer biases the solution toward identity scaling and limits the otherwise unbounded scale gauge. We use two variants. *Joint permutation+scale* optimizes Z_ℓ and u_ℓ together. *Scale refinement* first obtains hard permutations, freezes them, and optimizes only u_ℓ . The second variant can be applied after any permutation method and, when initialized with an exhaustive oracle, cleanly attributes any improvement to scale rather than a changed correspondence.

3.1 WHY AN EXACT SYMMETRY CHANGES THE INTERPOLATION PATH

Although scaling leaves the transformed endpoint function unchanged, it does not leave the functions in the interior of a straight parameter-space path unchanged. This distinction is the source of its utility. Consider one hidden ReLU unit with incoming weight w and outgoing weight v , and suppose for illustration that the activation pattern stays fixed along the segment. At the midpoint between unit (w_A, v_A) and a scaled representative $(dw_B, v_B/d)$, its locally linear contribution contains

$$\frac{1}{4} [v_A w_A + v_B w_B + d v_A w_B + d^{-1} v_B w_A]^\top x. \quad (7)$$

The first two terms are determined by the endpoints, but the two cross terms depend on d . Positive scaling can therefore balance the incompatible incoming and outgoing features that are mixed by interpolation. With multiple units and changing ReLU gates the expression is piecewise rather than globally linear, but the same mechanism remains: choosing a different representative on the symmetry orbit changes cross-model interactions while preserving both endpoint functions.

This also clarifies the different roles of width and symmetry. For a fixed width, permutation provides a finite orbit and scaling adds continuous directions within each permutation cell. Increasing width does more than enlarge that orbit: it changes the function class and can distribute a feature across redundant units, making both linear matching and curved motion easier. The experiments below are designed to observe these effects separately rather than treating every improvement as a better permutation search.

4 EXPERIMENTAL DESIGN

XOR endpoints. We train one-hidden-layer ReLU MLPs on the four XOR inputs with binary labels, using full-batch SGD and binary cross-entropy. Models are initialized uniformly in $[-0.5, 0.5]$, trained for at most 5000 epochs, and retained only when they reach 100% accuracy and low endpoint loss. For width h , all $h!$ hidden unit permutations are enumerated. The best permutation minimizes the sampled interpolation barrier, providing an exact oracle within permutation symmetry. For scale-aware experiments, interpolation is evaluated at 61 values of t . We obtain 21, 36, and 10 valid pairs at $h = 3, 5, 7$.

Nonlinear paths. To distinguish linear alignment from broader connectivity, we optimize both quadratic Bézier paths and polygonal chains between independently trained XOR endpoints. We begin with one trainable interior point and increase path flexibility up to ten interior points. The objective is the maximum loss over 31 sampled path locations.

Training-stage benchmark. For VGG16 on CIFAR-10, two branches share an initialization and the first n of $N = 200$ epochs, then train independently with different minibatch orders and augmentations. Later splits produce more closely related endpoints. A known random channel permutation is applied to one endpoint, and weight matching is asked to undo the mismatch. Because applying the inverse known permutation recovers the original branch, each split has a reference alignment. Repeating weight matching with different layer-update orders measures search variability.

VGG validation. We evaluate fixed independently trained seed-0/seed-1 pairs of VGG11, VGG13, VGG16, and VGG19 on CIFAR-10. We compare raw interpolation, permutation-only Sinkhorn, joint permutation+scale, and permutation followed by scale refinement. Alignment uses a deterministic calibration batch and minimizes midpoint loss; final results use hard permutations and test-set interpolation. Because each architecture is represented by one endpoint pair and method-wise hyperparameters are selected separately, this experiment is a realistic validation of the qualitative pattern, not a population estimate.

Additional optimization settings and reproducibility details appear in Appendix A.

5 PERMUTATION INSUFFICIENCY AND THE ROLE OF WIDTH


5.1 AN EXHAUSTIVE COUNTEREXAMPLE

At the minimal width $h = 2$, only the identity and swap are valid permutations. Several independently trained pairs retain a substantial barrier under both. Figure 1 shows representative cases. Alignment succeeds for seeds 2 and 4 because their parameters are nearly identical after a swap. In contrast, seeds 9 and 12 implement the same XOR labels and have visually similar decision boundaries, yet their internal weights are not related by either permutation and the barrier remains. A second failure mode occurs when perfect-accuracy solutions implement different boundaries away from the four training inputs. Since the search enumerates the full group, these failures cannot be blamed on the permutation optimizer. They form a concrete counterexample to a universal permutation-only LMC claim for trained networks; they do not contradict a claim restricted to sufficiently wide networks or to “most” solutions.

5.2 WIDTH LOWERS THE ORACLE BARRIER

Increasing width changes the result. In the initial exhaustive sweep over $h \in \{2, 3, 4, 5\}$, the number of pairs maintaining 100% accuracy along the best-permuted interpolation rises steadily; at $h = 5$, every evaluated pair maintains perfect accuracy. Perfect accuracy does not mean zero loss barrier: midpoint confidence can still fall relative to the endpoints.

The larger scale-aware study quantifies this effect in loss. Table 1 shows that the oracle permutation barrier falls from 0.357 at $h = 3$ to 0.111 at $h = 5$ and 0.023 at $h = 7$. Thus, even before scaling, width supplies matching flexibility. Practical Sinkhorn permutation search is weaker than the oracle at all widths, illustrating why heuristic failure alone cannot diagnose the absence of an alignment.



(a) Successful pair: a swap makes the endpoints nearly identical. (b) Failed pair: similar boundary, but neither permutation removes the barrier.

Figure 1: Exhaustive permutation alignment for width-2 XOR networks. Columns in each panel are interpolation points; top and bottom rows show raw and best-permuted interpolation. Both endpoints classify all four XOR points correctly.

Table 1: XOR loss barrier (mean $\pm$ standard deviation). Oracle+scale first selects the exhaustive best permutation, then optimizes only positive scales. The final row counts pairs for which refinement is no worse than the oracle.

Method	$h = 3$	$h = 5$	$h = 7$
No alignment	0.780 ± 0.610	0.608 ± 0.465	0.331 ± 0.184
Sinkhorn permutation	0.444 ± 0.325	0.189 ± 0.162	0.063 ± 0.044
Oracle permutation	0.357 ± 0.309	0.111 ± 0.118	0.023 ± 0.015
Joint permutation+scale	0.428 ± 0.380	0.074 ± 0.149	0.009 ± 0.009
Oracle+scale	0.247 ± 0.283	0.025 ± 0.044	0.005 ± 0.003
Oracle+scale $\leq$ oracle	21/21	36/36	10/10

5.3 NONLINEAR CONNECTIVITY CHANGES AT WIDTH THREE

Permutation-based LMC and general mode connectivity need not change at the same width. At $h = 2$, extensive Bézier and polygonal searches find low-loss paths only for essentially identical endpoints; additional bends do not recover nontrivial connections. At $h = 3$, all studied endpoint pairs are connected by low-loss paths. Most are found directly; for one difficult pair, low-loss paths through a third trained mode establish connectivity by concatenation. Figure 2 contrasts a width-2 failure with a width-3 success between models having different decision boundaries.

This is evidence about the paths found by the tested parameterizations, not a topological impossibility theorem for width two. Nevertheless, the sharp empirical transition is informative: adding a single hidden unit provides representational slack that permits a curved route even when the best straight symmetry-aligned route can retain loss. Width therefore affects connectivity through more than an enlarged permutation group.

6 SCALE IMPROVES ON THE PERMUTATION ORACLE

The decisive comparison in Table 1 is between oracle permutation and oracle+scale. The hard assignment is identical in the two conditions, so the improvements can only come from selecting a different functionally equivalent scaling of the same endpoint. Refinement reduces the mean barrier by 31% at $h = 3$, 77% at $h = 5$, and 78% at $h = 7$. It is no worse on every one of the 67 pairs. This shows that positive scale is not merely helping a relaxed optimizer discover a better permutation; it enlarges the attainable set beyond the exhaustive permutation orbit.

Joint optimization is more mixed at the narrowest width. It improves over the permutation oracle on 9/21, 29/36, and 8/10 pairs at widths 3, 5, and 7, respectively. At $h = 3$, its mean barrier (0.428) remains above the oracle-permutation result (0.357), whereas at widths five and seven it beats the oracle mean. The difference between joint and oracle-refinement results exposes an optimization gap: a richer symmetry class can contain better representatives without a from-scratch solver reliably finding them.




Figure 2: Nonlinear XOR paths. *Top*: a representative width-2 pair for which optimizing a Bézier path does not remove the high-loss region. *Bottom*: at width 3, a low-loss path connects endpoints with different decision boundaries. The rows visualize predictions along the learned path.




Figure 3: Mean XOR interpolation loss at widths 3, 5, and 7: raw interpolation (gray), Sinkhorn permutation (orange), exhaustive oracle permutation (blue), joint permutation+scale (purple), and oracle permutation followed by scale refinement (green). Exact barriers are in Table 1.

Figure 3 shows the averaged interpolation profiles. All curves improve with width, but oracle scale refinement (green) is consistently lowest. At $h = 7$, its mean maximum excess is only 0.005, compared with 0.023 for the best permutation. Scaling matters most once width has already made useful correspondences available: it fine-tunes the balance of incoming and outgoing weights without changing endpoint predictions.

7 WHEN DOES ALIGNMENT BECOME DIFFICULT?

The shared-training construction provides a dynamic view. Two VGG16 branches are identical until the split, after which SGD noise makes them diverge. A random known permutation is then added to one endpoint. Figure 4 plots test-error barrier against endpoint distance before and after weight matching. Late splits lie in the lower-left: their endpoints remain close, the reference path already has low barrier, and weight matching reliably removes the artificial mismatch. Early splits lie farther away and retain much larger barriers.

For every shared-split pair, the recovered alignment is at least as good as the known inverse permutation and often better. Repeating weight matching with different layer-update orders produces small variability (the largest barrier standard deviation is about 0.68 percentage points for the 0/200 split and becomes negligible for late splits). The trend is therefore not explained by one unlucky search order.

Distance is informative but not sufficient. The 4/196 shared-split pair has a test-error barrier of approximately 21.7 percentage points near distance 28, whereas the independently initialized pair after matching has a barrier of 34.1 points at a comparable distance. The experiment suggests




Figure 4: VGG16/CIFAR-10 shared-training benchmark. Labels $n/(200 - n)$ denote shared/independent epochs. Circles are original endpoints; squares are endpoints after applying a random permutation and weight matching. Red denotes independent initialization. Later splits remain closer and have lower barriers after alignment.

that weight similarity is a useful proxy when trajectories share substantial history, but ceases to characterize linear compatibility once their learned representations diverge. This motivates optimizing interpolation loss within a broader symmetry orbit, as in Section 3.

8 VGG/CIFAR-10 VALIDATION

Figure 5 reports test-loss barriers for four VGG architectures. Permutation-only Sinkhorn reduces the raw barriers from roughly 1.88–1.98 to 0.38, 0.79, 1.22, and 1.21 for VGG11/13/16/19. Adding scales improves on permutation-only in every architecture. The better of joint and sequential scaling attains 0.109, 0.429, 1.035, and 1.105, respectively. Relative to permutation-only, these are reductions of about 71%, 45%, 15%, and 9%.

Neither optimization order dominates. Joint permutation+scale is better for VGG11 and VGG19; freezing the permutation and refining scales is better for VGG13 and VGG16. This echoes XOR: the symmetry class contains better representatives, but joint nonconvex optimization does not always realize the full gain. Substantial barriers also remain for deeper VGGs, so permutation and scale together are not a complete account of LMC.

The VGG results should be read as a transfer check of the controlled XOR finding. Exhaustive permutation search is impossible, only one endpoint pair is used per architecture, and hyperparameters are selected per method. Nevertheless, the consistent direction of change across four architectures supports the practical relevance of scaling beyond the toy setting.

9 DISCUSSION AND LIMITATIONS

A refined one-basin picture. Our experiments separate three claims that are easy to collapse into one. First, curved connectivity can exist without a good straight path. Second, permutations can remove representational mismatch, but the exhaustive XOR result shows that this exact symmetry




Figure 5: Maximum test-loss barrier for independently trained VGG models on CIFAR-10. Both scale-aware variants improve on Sinkhorn permutation-only for all four fixed endpoint pairs, although joint and sequential optimization trade places across architectures.

is not always sufficient. Third, overparameterization can make both processes easier by supplying redundant features and routes through function space. Positive scale adds another exact degree of freedom and systematically improves the oracle permutation, but does not eliminate every barrier. The evidence therefore favors connectivity modulo a richer symmetry class over a literal permutation-only one-basin claim.

Optimization versus existence. The exhaustive protocol is central to this interpretation. Without it, a high barrier after Sinkhorn or weight matching could always be ascribed to search failure. Conversely, the gap between joint Sinkhorn and oracle refinement warns against treating a failed continuous relaxation as evidence that a good monomial representative does not exist. Future work should distinguish the expressivity of a transformation class from the solver used to search it.

A diagnostic hierarchy for alignment failures. The experiments suggest a practical order for interpreting a residual barrier. First, establish whether search error can be removed. Exhaustive enumeration is available only for small networks, but multiple restarts, exact matching in restricted layers, or controlled known-permutation benchmarks can still reveal optimizer limitations. Second, if the best permutation remains inadequate, expand the exact symmetry class before concluding that modes occupy distinct basins. Scale refinement is useful here because it is continuous, function preserving after the permutation is fixed, and can be attached to existing aligners. Third, compare linear alignment with nonlinear paths and width changes. A curved path that appears after adding a unit is evidence for representational slack, not evidence that a hidden permutation was recovered. This hierarchy separates what transformation exists, whether an algorithm finds it, and whether the architecture has enough redundancy for the desired path.

For model merging, the results also favor a direct validation rule: endpoint equivalence should be checked after every transformation, and the full interpolation profile should be measured on held-out data. Small Euclidean distance is not by itself a certificate of compatibility, as the training-stage experiment demonstrates. Conversely, a scale-aware transformation can redistribute parameter differences while lowering the loss barrier. Objectives designed for merging should therefore target functional behavior along the path, using weight distance only as a computational proxy.

Limitations. The exact conclusions come from small one-hidden-layer XOR networks. The VGG study is non-exhaustive, uses a single pair per architecture, and relies on method-specific hyperparameter selection. Our nonlinear path searches can certify found paths but cannot certify nonexistence. The scale-aware soft transformation is exactly function preserving only after projection; during Sinkhorn optimization it is a relaxation. Finally, we consider ReLU networks, for which positive homogeneity holds, and do not address normalization layers, residual architectures, attention symmetries, or approximate functional equivalences. Larger multi-seed studies and comparisons with activation repair are needed before drawing broad architectural conclusions.

10 CONCLUSION

Permutation symmetry is an important but incomplete explanation of linear mode connectivity. Exhaustive XOR search exposes genuine permutation-only failures; increasing width rapidly suppresses them and changes nonlinear connectivity; and positive scaling improves even the best possible permutation. The same directional gain appears in VGG/CIFAR-10, while a shared-training experiment shows how alignment difficulty grows as optimization trajectories diverge. Together, the results suggest that the useful object is not a basin modulo permutation alone, but a low-loss geometry shaped jointly by exact symmetries, width, and the history of optimization.

AI USE STATEMENT

Generative AI tools were used to discuss ideas, improve the organization and clarity of draft text, and assist with programming tasks including debugging, refactoring, plotting, and $\LaTeX$ editing. The research questions, experimental design, implementation decisions, analysis, and conclusions were determined by the authors. AI-assisted text and code were reviewed and checked before inclusion. The authors take responsibility for the final content, claims, citations, and artifacts. This statement should be reconciled with the final author team’s use before submission.

REFERENCES

- Samuel Ainsworth, Jonathan Hayase, and Siddhartha Srinivasa. Git re-basin: Merging models modulo permutation symmetries. In *The Eleventh International Conference on Learning Representations*, 2023. URL <https://openreview.net/forum?id=CQsmMYmlP5T>.
- Felix Draxler, Kambis Veschgini, Manfred Salmhofer, and Fred Hamprecht. Essentially no barriers in neural network energy landscape. In Jennifer Dy and Andreas Krause (eds.), *Proceedings of the 35th International Conference on Machine Learning*, volume 80 of *Proceedings of Machine Learning Research*, pp. 1309–1318. PMLR, 10–15 Jul 2018. URL <https://proceedings.mlr.press/v80/draxler18a.html>.
- Rahim Entezari, Hanie Sedghi, Olga Saukh, and Behnam Neyshabur. The role of permutation invariance in linear mode connectivity of neural networks. In *International Conference on Learning Representations*, 2021. URL <https://openreview.net/forum?id=dNigytemkL>.
- Jonathan Frankle, Gintare Karolina Dziugaite, Daniel Roy, and Michael Carbin. Linear mode connectivity and the lottery ticket hypothesis. In Hal Daumé III and Aarti Singh (eds.), *Proceedings of the 37th International Conference on Machine Learning*, volume 119 of *Proceedings of Machine Learning Research*, pp. 3259–3269. PMLR, 13–18 Jul 2020. URL <https://proceedings.mlr.press/v119/frankle20a.html>.
- Timur Garipov, Pavel Izmailov, Dmitrii Podoprikin, Dmitry P. Vetrov, and Andrew G. Wilson. Loss surfaces, mode connectivity, and fast ensembling of DNNs. In *Advances in Neural Information Processing Systems*, volume 31, 2018. URL <https://papers.nips.cc/paper/8095-loss-surfaces-mode-connectivity-and-fast-ensembling-of-dnns>.
- Pavel Izmailov, Dmitrii Podoprikin, Timur Garipov, Dmitry Vetrov, and Andrew Gordon Wilson. Averaging weights leads to wider optima and better generalization, 2019. URL <https://arxiv.org/abs/1803.05407>.

Keller Jordan, Hanie Sedghi, Olga Saukh, Rahim Entezari, and Behnam Neyshabur. REPAIR: RENormalizing permuted activations for interpolation repair. In *International Conference on Learning Representations*, 2023. URL <https://openreview.net/forum?id=gU5sJ6ZggcX>.

Fidel A. Guerrero Peña, Heitor Rapela Medeiros, Thomas Dubail, Masih Aminbeidokhti, Eric Granger, and Marco Pedersoli. Re-basin via implicit sinkhorn differentiation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 20237–20246, June 2023. URL https://openaccess.thecvf.com/content/CVPR2023/html/Pena_Re-Basin_via_Implicit_Sinkhorn_Differentiation_CVPR_2023_paper.html.

Mitchell Wortsman, Gabriel Ilharco, Samir Ya Gadre, Rebecca Roelofs, Raphael Gontijo-Lopes, Ari S Morcos, Hongseok Namkoong, Ali Farhadi, Yair Carmon, Simon Kornblith, and Ludwig Schmidt. Model soups: averaging weights of multiple fine-tuned models improves accuracy without increasing inference time. In Kamalika Chaudhuri, Stefanie Jegelka, Le Song, Csaba Szepesvari, Gang Niu, and Sivan Sabato (eds.), *Proceedings of the 39th International Conference on Machine Learning*, volume 162 of *Proceedings of Machine Learning Research*, pp. 23965–23998. PMLR, 17–23 Jul 2022. URL <https://proceedings.mlr.press/v162/wortsman22a.html>.

A REPRODUCIBILITY DETAILS

A.1 XOR TRAINING AND EXHAUSTIVE ALIGNMENT

The XOR set is $\{(0, 0), (0, 1), (1, 0), (1, 1)\}$ with labels $\{0, 1, 1, 0\}$. Models have one ReLU hidden layer, biases, and one output logit. Full-batch SGD uses learning rate 0.05 and at most 5000 epochs. ReduceLROnPlateau uses factor 0.5, patience 200, threshold 10^{-4} , and minimum learning rate 5×10^{-4} . Training stops when loss has not improved by 10^{-4} for 500 epochs. For each retained pair, all $h!$ permutations are evaluated. Interpolation selection prioritizes classification error in the width-2 diagnostic and loss barrier in the scale-aware study.

For nonlinear paths, both quadratic Bézier curves and polygonal chains are optimized for 1500 steps at learning rate 0.05, with no weight decay and 31 path samples. The default path has one trainable interior point; harder pairs use up to ten.

A.2 XOR SCALE-AWARE OPTIMIZATION

Final interpolation curves use 61 evaluation points; optimization uses 31 points. The default permutation-only Sinkhorn run uses Adam for 500 steps, learning rate 0.01, temperature 1, 20 Sinkhorn iterations, and identity strength 1. Scale refinement uses Adam for 500 steps, learning rate 0.01, and log-scale regularization 10^{-3} . Joint optimization uses the same default budget. If a default run misses its target, a bounded search over learning rates, temperatures, and regularization is used; settings are selected per pair. The scale target is barrier $\epsilon = 0.01$, while permutation-only Sinkhorn targets that pair’s exhaustive barrier.

A.3 VGG EXPERIMENTS

The shared-training benchmark follows a 200-epoch VGG16/CIFAR-10 schedule. Branches differ in minibatch order and stochastic augmentation after the split. The artificial full-network permutation uses seed 42; weight matching uses at most 100 block-coordinate iterations and is repeated with three layer-update orders.

For the scale-aware comparison, fixed seed-0/seed-1 endpoint pairs are used for VGG11/13/16/19. Alignment uses a deterministic batch of 1000 images, 20 Sinkhorn iterations, and midpoint loss. Separate bounded sweeps select the number of optimization steps, learning rate, temperature, and scale regularization for each architecture and method. The permutation-then-scale baseline freezes the projected hard permutation and optimizes only log-scales.

B ADDITIONAL PAIRWISE RESULTS

Table 2 separates the practical optimization gap from the gain provided by the symmetry class. Sinkhorn permutation-only exactly matches the exhaustive oracle on progressively fewer pairs as width grows. Joint permutation+scale often beats the permutation oracle, but oracle scale refinement is the most reliable condition.

Table 2: Pairwise XOR comparisons. Lower loss barrier is better.

Comparison	$h = 3$	$h = 5$	$h = 7$
Oracle+scale $\leq$ oracle permutation	21/21	36/36	10/10
Joint permutation+scale $\leq$ oracle permutation	9/21	29/36	8/10
Sinkhorn permutation = oracle permutation	9/21	7/36	0/10
Oracle+scale $\leq$ joint permutation+scale	15/21	25/36	5/10